

In [1]:

```
1 # ML Facial recognition to detect mood and suggest songs accordingly
```

In [2]:

```
1 import numpy as np
2 import pandas as pd
```

C:\Users\tridi\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.7.3' or newer of 'numexpr' (version '2.7.1' currently installed).  
from pandas.core.computation.check import NUMEXPR\_INSTALLED

In [3]:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
```

In [4]:

```
1 import cv2
2 import os
3 import random
4 from skimage.io import imread
```

In [5]:

```
1 import tensorflow as tf
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense, Dropout, Flatten, BatchNormalization, Conv2D, MaxPooling2D, Activation
4 from tensorflow.keras.optimizers import Adam
5 from tensorflow.keras.preprocessing.image import ImageDataGenerator
6 from tensorflow.keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
7 from tensorflow.keras.models import load_model
```

In [6]:

```
1 from IPython.display import Audio
```

In [7]:

```
1 # read dataset
2
3 from skimage.io import imread
4 train_dir = "dataset/train/"
5 test_dir = "dataset/test/"
```

In [8]:


```
1 total_labels = len(os.listdir(train_dir))
```

In [9]:

```

1 #data visualizatin
2 fig, ax = plt.subplots(nrows=5, ncols=total_labels, figsize=(35, 25))
3 for x in range(5):
4     for y,v in zip(range(total_labels),os.listdir(train_dir)):
5         ax[x][y].imshow(imread(train_dir+v+'/' +os.listdir(train_dir+v)[x]))
6
7 plt.show()

```



In [10]:

```

1 #gray images for cnn
2 fig, ax = plt.subplots(nrows=5, ncols=total_labels, figsize=(35, 25))
3 for x in range(5):
4     for y,v in zip(range(total_labels),os.listdir(train_dir)):
5         ax[x][y].imshow(imread(train_dir+v+'/' +os.listdir(train_dir+v)[x]), cmap='gray')
6
7 plt.show()

```

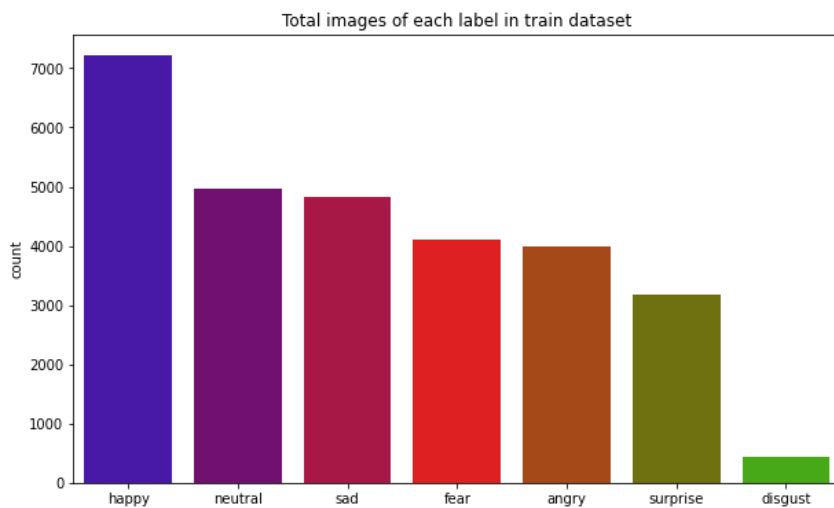


In [11]:

```

1 df = {}
2 for i in os.listdir(train_dir):
3     directory = train_dir + i
4     df[i] = len(os.listdir(directory))
5 df = pd.DataFrame(df, index=["total"]).transpose().sort_values("total", ascending=False)
6
7 plt.figure(figsize=(10,6))
8 sns.barplot(x=df.index, y="total", palette="brg", data=df)
9 plt.ylabel("count")
10 plt.title("Total images of each label in train dataset")
11 plt.show()

```



In [12]:

```

1 happy = os.listdir(train_dir+'happy/')
2 dim1, dim2 = [], []
3
4 for img_filename in happy:
5     img = imread(train_dir+'happy/'+img_filename)
6     d1, d2 = img.shape
7     dim1.append(d1)
8     dim2.append(d2)


```

In [13]:

```

1 img_shape = (int(np.mean(dim1)), int(np.mean(dim2)), 1)
2 sns.jointplot(dim1, dim2,color='red')
3 plt.show()

```



In [14]:

```
1 #Data Preprocessing
```

In [15]:

```

1 #to generate batches of tensor images
2
3 train_gen = ImageDataGenerator(rescale=1/255,
4                               rotation_range=40,
5                               width_shift_range=0.2,
6                               height_shift_range=0.2,
7                               shear_range=0.2,
8                               zoom_range=0.2,
9                               horizontal_flip=True,
10                              fill_mode='nearest')
11
12 test_gen = ImageDataGenerator(rescale=1/255)

```

In [16]:

```
1 img_shape = (int(np.mean(dim1)), int(np.mean(dim2)), 1)
```

In [17]:

```

1 #splitting into train& test data
2
3 train_generator = train_gen.flow_from_directory(directory=train_dir,
4                                                target_size=(img_shape[0], img_shape[1]),
5                                                color_mode='grayscale',
6                                                batch_size=64,
7                                                class_mode='categorical',
8                                                shuffle=True)
9
10 test_generator = test_gen.flow_from_directory(directory=test_dir,
11                                              target_size=(img_shape[0], img_shape[1]),
12                                              color_mode='grayscale',
13                                              batch_size=64,
14                                              class_mode='categorical',
15                                              shuffle=False)

```

Found 28709 images belonging to 7 classes.  
Found 7178 images belonging to 7 classes.

In [18]:

```
1 # Creating the Model
```

In [19]:

```
1 model=Sequential()
```

In [20]:

```
1 # with 64 filter
2 model.add(Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu', input_shape=img_shape))
3 model.add(BatchNormalization())
4 model.add(MaxPooling2D(pool_size=(2,2)))
5 model.add(Dropout(0.2))
```

In [21]:

```
1 #with 128 filters
2 model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'))
3 model.add(BatchNormalization())
4 model.add(MaxPooling2D(pool_size=(2,2)))
5 model.add(Dropout(0.2))
```

In [22]:

```
1 # with 512
2 model.add(Conv2D(filters=512, kernel_size=(3,3), padding='same', activation='relu'))
3 model.add(BatchNormalization())
4 model.add(MaxPooling2D(pool_size=(2,2)))
5 model.add(Dropout(0.2))
```

In [23]:

```
1 model.add(Conv2D(filters=512, kernel_size=(3,3), padding='same', activation='relu'))
2 model.add(BatchNormalization())
3 model.add(MaxPooling2D(pool_size=(2,2)))
4 model.add(Dropout(0.2))
```

In [24]:

```
1 model.add(Flatten())
```

In [25]:

```
1 model.add(Dense(512,activation='relu'))
2 model.add(Dropout(0.2))
```

In [26]:

```
1 model.add(Dense(1024,activation='relu'))
2 model.add(Dropout(0.2))
```

In [27]:

```
1 model.add(Dense(units=len(os.listdir(train_dir)),activation='softmax'))
```

In [28]:

```
1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 48, 48, 64)	640
batch_normalization (Batch Normalization)	(None, 48, 48, 64)	256
max_pooling2d (MaxPooling2D)	(None, 24, 24, 64)	0
dropout (Dropout)	(None, 24, 24, 64)	0
conv2d_1 (Conv2D)	(None, 24, 24, 128)	73856
batch_normalization_1 (Batch Normalization)	(None, 24, 24, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_1 (Dropout)	(None, 12, 12, 128)	0
conv2d_2 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_2 (Batch Normalization)	(None, 12, 12, 512)	2048
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 512)	0
dropout_2 (Dropout)	(None, 6, 6, 512)	0
conv2d_3 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_3 (Batch Normalization)	(None, 6, 6, 512)	2048
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 512)	0
dropout_3 (Dropout)	(None, 3, 3, 512)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 512)	2359808
dropout_4 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 1024)	525312
dropout_5 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 7)	7175
=====		
Total params: 5,921,799		
Trainable params: 5,919,367		
Non-trainable params: 2,432		

In [29]:

```
1 # Train the model
```

In [30]:

```
1 model.compile(optimizer=Adam(learning_rate=0.0001,
2                             decay=1e-6),
3               loss='categorical_crossentropy',
4               metrics=['accuracy'])
```

In [31]:

```

1 steps_per_epoch = train_generator.n // train_generator.batch_size
2 validation_steps = test_generator.n // test_generator.batch_size
3 num_epochs = 13
4
5 history = model.fit(train_generator,
6                     epochs=num_epochs,
7                     verbose=1,
8                     #callbacks=callbacks,
9                     validation_data=test_generator,
10                    steps_per_epoch=steps_per_epoch,
11                    validation_steps=validation_steps)

```

```

Epoch 1/13
448/448 [=====] - 921s 2s/step - loss: 1.8810 - accuracy: 0.2267 - val_loss: 2.3236 - val_accu
racy: 0.1731
Epoch 2/13
448/448 [=====] - 779s 2s/step - loss: 1.8036 - accuracy: 0.2445 - val_loss: 1.7933 - val_accu
racy: 0.2662
Epoch 3/13
448/448 [=====] - 797s 2s/step - loss: 1.7865 - accuracy: 0.2563 - val_loss: 1.7478 - val_accu
racy: 0.2857
Epoch 4/13
448/448 [=====] - 830s 2s/step - loss: 1.7730 - accuracy: 0.2681 - val_loss: 1.7122 - val_accu
racy: 0.3086
Epoch 5/13
448/448 [=====] - 649s 1s/step - loss: 1.7537 - accuracy: 0.2805 - val_loss: 1.7661 - val_accu
racy: 0.3104
Epoch 6/13
448/448 [=====] - 616s 1s/step - loss: 1.7314 - accuracy: 0.2934 - val_loss: 1.7054 - val_accu
racy: 0.3340
Epoch 7/13
448/448 [=====] - 778s 2s/step - loss: 1.7114 - accuracy: 0.3016 - val_loss: 1.7847 - val_accu
racy: 0.3165
Epoch 8/13
448/448 [=====] - 855s 2s/step - loss: 1.6775 - accuracy: 0.3259 - val_loss: 1.5881 - val_accu
racy: 0.3850
Epoch 9/13
448/448 [=====] - 549s 1s/step - loss: 1.6370 - accuracy: 0.3444 - val_loss: 1.7637 - val_accu
racy: 0.3676
Epoch 10/13
448/448 [=====] - 552s 1s/step - loss: 1.6062 - accuracy: 0.3642 - val_loss: 1.4751 - val_accu
racy: 0.4067
Epoch 11/13
448/448 [=====] - 589s 1s/step - loss: 1.5656 - accuracy: 0.3857 - val_loss: 1.4236 - val_accu
racy: 0.4634
Epoch 12/13
448/448 [=====] - 589s 1s/step - loss: 1.5330 - accuracy: 0.3983 - val_loss: 1.3380 - val_accu
racy: 0.4895
Epoch 13/13
448/448 [=====] - 563s 1s/step - loss: 1.5031 - accuracy: 0.4165 - val_loss: 1.3777 - val_accu
racy: 0.4756

```

In [32]:

```
1 model.save('model.h5')
```

In [33]:

```

1 # Evaluate the model
2 test_loss, test_acc = model.evaluate(test_generator)
3 print("validation accuracy :", str(test_acc*100)+"%")
4 print("validation loss :", test_loss)

```

```

113/113 [=====] - 30s 268ms/step - loss: 1.3765 - accuracy: 0.4762
validation accuracy : 47.61772155761719%
validation loss : 1.3765020370483398

```

In [34]:

```
1 # Plotting Training and Validation plot
```


In [36]:

```

1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs = range(len(acc))
6
7 fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))
8 ax[0].plot(epochs, acc, 'red', label='Training accuracy')
9 ax[0].plot(epochs, val_acc, 'green', label='Validation accuracy')
10 ax[0].legend(loc=0)
11 ax[1].plot(epochs, loss, 'red', label='Training loss')
12 ax[1].plot(epochs, val_loss, 'green', label='Validation loss')
13 ax[1].legend(loc=0)
14
15 plt.suptitle('Training and validation')
16 plt.show()

```

Training and validation



In [41]:

```
1 # Testing our model with new image
```

In [ ]:

```
1 #1
```

In [38]:

```

1 image = cv2.imread("images/sad.jpg")
2 from IPython.display import Image
3 Image(filename='images/sad.jpg')

```

Out[38]:





In [40]:

```
1 # Model Prediction 1
2
3 model = load_model("model.h5")
4 # A List of emoticon categories
5 EMOTIONS = ['Angry', 'Disgust', 'Happy', 'Sad', 'Surprise', 'Neutral']
6 # Load image
7 img = image
8
9 # Trim the image to 48 x 48, and turn the grayscale image, normalization
10 frame = cv2.resize(img,(48,48),interpolation=cv2.INTER_BITS2)
11 gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY) / 255.0
12
13 # Reinvent the image dimension
14 gray = gray.reshape(1,48,48,1)
15
16 # Output the prediction
17 predicts = model.predict(gray)[0]
18 label = EMOTIONS[predicts.argmax()]
19 for (i,j) in zip(range(7),EMOTIONS):
20     predictss = predicts[i]
21     print("{:^10s}".format(j)+"prediction rate is  {0:.2f}%".format(predictss))
22 print( "\n\n The system considers this expression to be:",label)
```

```
1/1 [=====] - 0s 141ms/step
Angry    prediction rate is  0.22%
Disgust  prediction rate is  0.02%
Happy    prediction rate is  0.11%
Sad      prediction rate is  0.32%
Surprise prediction rate is  0.10%
Neutral  prediction rate is  0.18%
```

The system considers this expression to be: Sad

In [43]:

```
1 # Song Recommendation 1
2
3
4 if (label=='Angry'):
5     path="song\\Angry\\"
6     files=os.listdir(path)
7     d=random.choice(files)
8     print("Now Playing:",d)
9     audio = Audio(filename='song\\Angry\\'+ d,autoplay=True)
10    display(audio)
11
12 elif (label=='Disgust'):
13     path="song\\Disgust\\"
14     files=os.listdir(path)
15     d=random.choice(files)
16     print("Now Playing:",d)
17     audio = Audio(filename='song\\Disgust\\'+ d,autoplay=True)
18    display(audio)
19
20 elif (label=="Happy"):
21     path="song\\Happy\\"
22     files=os.listdir(path)
23     d=random.choice(files)
24     print("Now Playing:",d)
25     audio = Audio(filename='song\\Happy\\'+ d,autoplay=True)
26    display(audio)
27
28 elif (label=='Sad'):
29     path="song\\Sad\\"
30     files=os.listdir(path)
31     d=random.choice(files)
32     print("Now Playing:",d)
33     audio = Audio(filename='song\\Sad\\'+ d,autoplay=True)
34    display(audio)
35
36 elif (label=='Surprise'):
37     path="song\\Surprise\\"
38     files=os.listdir(path)
39     d=random.choice(files)
40     print("Now Playing:",d)
41     audio = Audio(filename='song\\Surprise\\'+ d,autoplay=True)
42    display(audio)
43
44 elif (label=='Neutral'):
45     path="song\\Neutral\\"
46     files=os.listdir(path)
47     d=random.choice(files)
48     print("Now Playing:",d)
49     audio = Audio(filename='song\\Neutral\\'+ d,autoplay=True)
50    display(audio)
```

Now Playing: Srotoshini - Encore.mp3

0:17 / 4:01

In [44]:

```
1 #2
```

In [48]:

```

1 image = cv2.imread("images/angry.jpeg")
2 from IPython.display import Image
3 Image(filename='images/angry.jpeg')

```

Out[48]:



In [49]:

```

1 # Model Prediction 2
2
3 model = load_model("model.h5")
4 # A List of emoticon categories
5 EMOTIONS = ['Angry', 'Disgust', 'Happy', 'Sad', 'Surprise', 'Neutral']
6 # Load image
7 img = image
8
9 # Trim the image to 48 x 48, and turn the grayscale image, normalization
10 frame = cv2.resize(img,(48,48),interpolation=cv2.INTER_BITS2)
11 gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY) / 255.0
12
13 # Reinvent the image dimension
14 gray = gray.reshape(1,48,48,1)
15
16 # Output the prediction
17 predicts = model.predict(gray)[0]
18 label = EMOTIONS[predicts.argmax()]
19 for (i,j) in zip(range(7),EMOTIONS):
20     predictss = predicts[i]
21     print("{:^10s}".format(j)+"prediction rate is  {0:.2f}%".format(predictss))
22 print( "\n\n The system considers this expression to be:",label)

```

```

1/1 [=====] - 0s 109ms/step
Angry    prediction rate is  0.24%
Disgust  prediction rate is  0.01%
Happy    prediction rate is  0.20%
Sad      prediction rate is  0.16%
Surprise prediction rate is  0.14%
Neutral  prediction rate is  0.17%

```

The system considers this expression to be: Angry

In [50]:

```
1 # Song Recommendation 2
2
3
4 if (label=='Angry'):
5     path="song\\Angry\\"
6     files=os.listdir(path)
7     d=random.choice(files)
8     print("Now Playing:",d)
9     audio = Audio(filename='song\\Angry\\'+ d,autoplay=True)
10    display(audio)
11
12 elif (label=='Disgust'):
13     path="song\\Disgust\\"
14     files=os.listdir(path)
15     d=random.choice(files)
16     print("Now Playing:",d)
17     audio = Audio(filename='song\\Disgust\\'+ d,autoplay=True)
18    display(audio)
19
20 elif (label=="Happy"):
21     path="song\\Happy\\"
22     files=os.listdir(path)
23     d=random.choice(files)
24     print("Now Playing:",d)
25     audio = Audio(filename='song\\Happy\\'+ d,autoplay=True)
26    display(audio)
27
28 elif (label=='Sad'):
29     path="song\\Sad\\"
30     files=os.listdir(path)
31     d=random.choice(files)
32     print("Now Playing:",d)
33     audio = Audio(filename='song\\Sad\\'+ d,autoplay=True)
34    display(audio)
35
36 elif (label=='Surprise'):
37     path="song\\Surprise\\"
38     files=os.listdir(path)
39     d=random.choice(files)
40     print("Now Playing:",d)
41     audio = Audio(filename='song\\Surprise\\'+ d,autoplay=True)
42    display(audio)
43
44 elif (label=='Neutral'):
45     path="song\\Neutral\\"
46     files=os.listdir(path)
47     d=random.choice(files)
48     print("Now Playing:",d)
49     audio = Audio(filename='song\\Neutral\\'+ d,autoplay=True)
50    display(audio)
```

Now Playing: Sadda Haq-rockstar.mp3

0:10 / 6:48

In [ ]:

```
1 # Thank You
```